



# Ten Ideas for webenar

Growth wizzards

August-31-2026

Yazan Alkhadraa

### **1. Designing a Smart LLM Router: How to Cut Model API Costs by 70%**

Learn how leading AI teams route requests across multiple LLMs to balance cost, speed, and accuracy. This session breaks down real routing strategies that can slash your inference bill without sacrificing output quality.

### **2. Common Mistakes When Building Your First AI Agent (and How to Avoid Them)**

A practical, no-fluff session covering the pitfalls most beginners hit when building AI agents — from poor tool design to runaway loops and unreliable outputs. Walk away with a checklist to build more robust agents from day one.

### **3. Fine-tuning vs. RAG vs. Prompt Engineering: When to Use Each**

Confused about which approach fits your use case? This session breaks down the trade-offs between fine-tuning, retrieval-augmented generation, and prompt engineering — with real examples of when each one wins.

### **4. Build a Smart Chatbot Connected to Your Own Database in One Hour**

A hands-on, live-build session where we create a chatbot that can query your own data in real time using RAG. Perfect for beginners who want a working project by the end of the hour.

### **5. From Idea to Intelligent Agent: Building an AI Agent in 60 Minutes**

A live, step-by-step build of a functional AI agent — covering planning, tool use, and execution. No prior agent-building experience required.

## **6. Train Your Own LLM: From Raw Data to Fine-Tuned Model**

An introduction to the end-to-end process of fine-tuning a language model — data preparation, training considerations, and evaluation. Ideal for those curious about what happens beyond calling an API.

## **7. AI Infrastructure Explained: Why Companies Choose Specific GPUs**

A behind-the-scenes look at the hardware powering modern AI systems — what makes a GPU suitable for training and inference, and how infrastructure choices affect cost and performance at scale.

## **8. From Zero to Your First AI Engineer Job: A Practical Roadmap**

A career-focused session mapping out the skills, projects, and learning path needed to break into AI engineering — with honest advice on what actually gets you hired.

## **9. Protecting Data in the Age of LLMs: An Introduction to PII Masking**

Learn why safeguarding personal data before it reaches an LLM is critical, and explore practical techniques for detecting and masking sensitive information in real-world pipelines.

## **10. Prompt Engineering vs. Context Engineering vs. Harness Engineering: What's the Difference?**

As AI systems grow more complex, so does the vocabulary around building them. This session clarifies these three related-but-distinct disciplines and shows how they fit together in real agent development.

## **11. What Are Neural Networks?**

A beginner-friendly, visual introduction to how neural networks actually work — no math background required. Perfect for anyone who's heard the term everywhere but wants to finally understand the fundamentals.